

XAAS: ACCELERATION AS A SERVICE TO ENABLE PRODUCTIVE HIGH-PERFORMANCE CLOUD COMPUTING


Panelists:

- Dan Ernst (NVIDIA)
- Ian Foster (Argonne National Laboratory)
- Torsten Hoefler (ETH Zurich)
- Thomas C. Schulthess (CSCS)

Moderated by Marcin Copik (ETH Zurich)



What is Axeleration as a Service?



Theme Article: Converged Computing: A Best-of-Both Worlds of HPC and Cloud

XaaS: Acceleration as a Service to Enable Productive High-Performance Cloud Computing

Torsten Hoefer, *ETH Zurich & Swiss National Supercomputing Centre (CSCS), Switzerland*

Marcin Copik, *ETH Zurich, Switzerland*

Pete Beckman, *Argonne National Laboratory, USA*

Andrew Jones, *Microsoft, United Kingdom*

Ian Foster, *Argonne National Laboratory, USA*

Manish Parashar, *Utah University, USA*

Daniel Reed, *Utah University, USA*

Matthias Troyer, *Microsoft, USA*

Thomas Schulthess, *Swiss National Supercomputing Centre (CSCS), Switzerland*

Dan Ernst, *NVIDIA, USA*

Jack Dongarra, *University of Tennessee, USA*




Paper

Performance Portable Containers

cloud computing



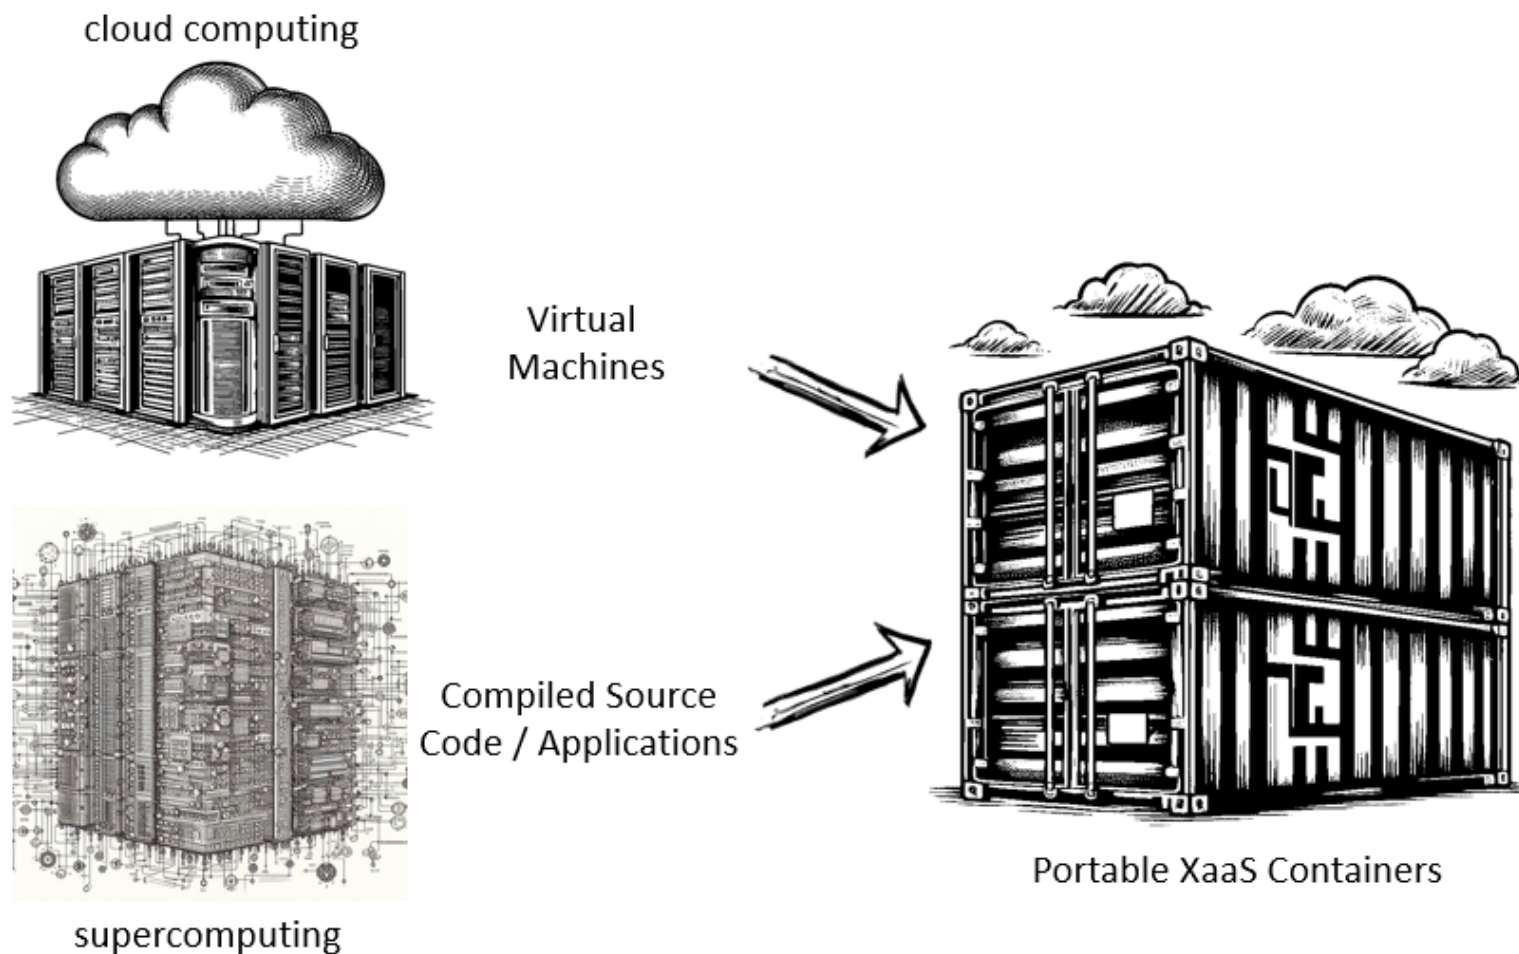
Virtual
Machines




Compiled Source
Code / Applications

supercomputing


Performance Portable Containers



Performance Portable Containers




Three Building Blocks of XaaS



Portable XaaS Containers

Three Building Blocks of XaaS



Portable XaaS Containers

High-Performance I/O

- Data storage managed by system provider
- Communication between XaaS containers
- Specialization to network fabric

Three Building Blocks of XaaS



Portable XaaS Containers

High-Performance I/O

- Data storage managed by system provider
- Communication between XaaS containers
- Specialization to network fabric


Scheduling & Invocations

- Flexible scheduling to reduce wait times
- Balance between batch and interactive workloads
- High-level control plane API, e.g., REST


XaaS System Architecture



XaaS System Architecture




XaaS System Architecture



XaaS System Architecture

XaaS Accelerated Compute, Communication, I/O libraries, and APIs



System-Specific
(e.g., HPE, Intel)




Base Operating System Layer (e.g., Ubuntu/Debian/... - POSIX)







Generic Container
(e.g., docker)



Standard Portable Container API
(e.g., OCI)

   ...

Commercial Cloud Providers

    ...

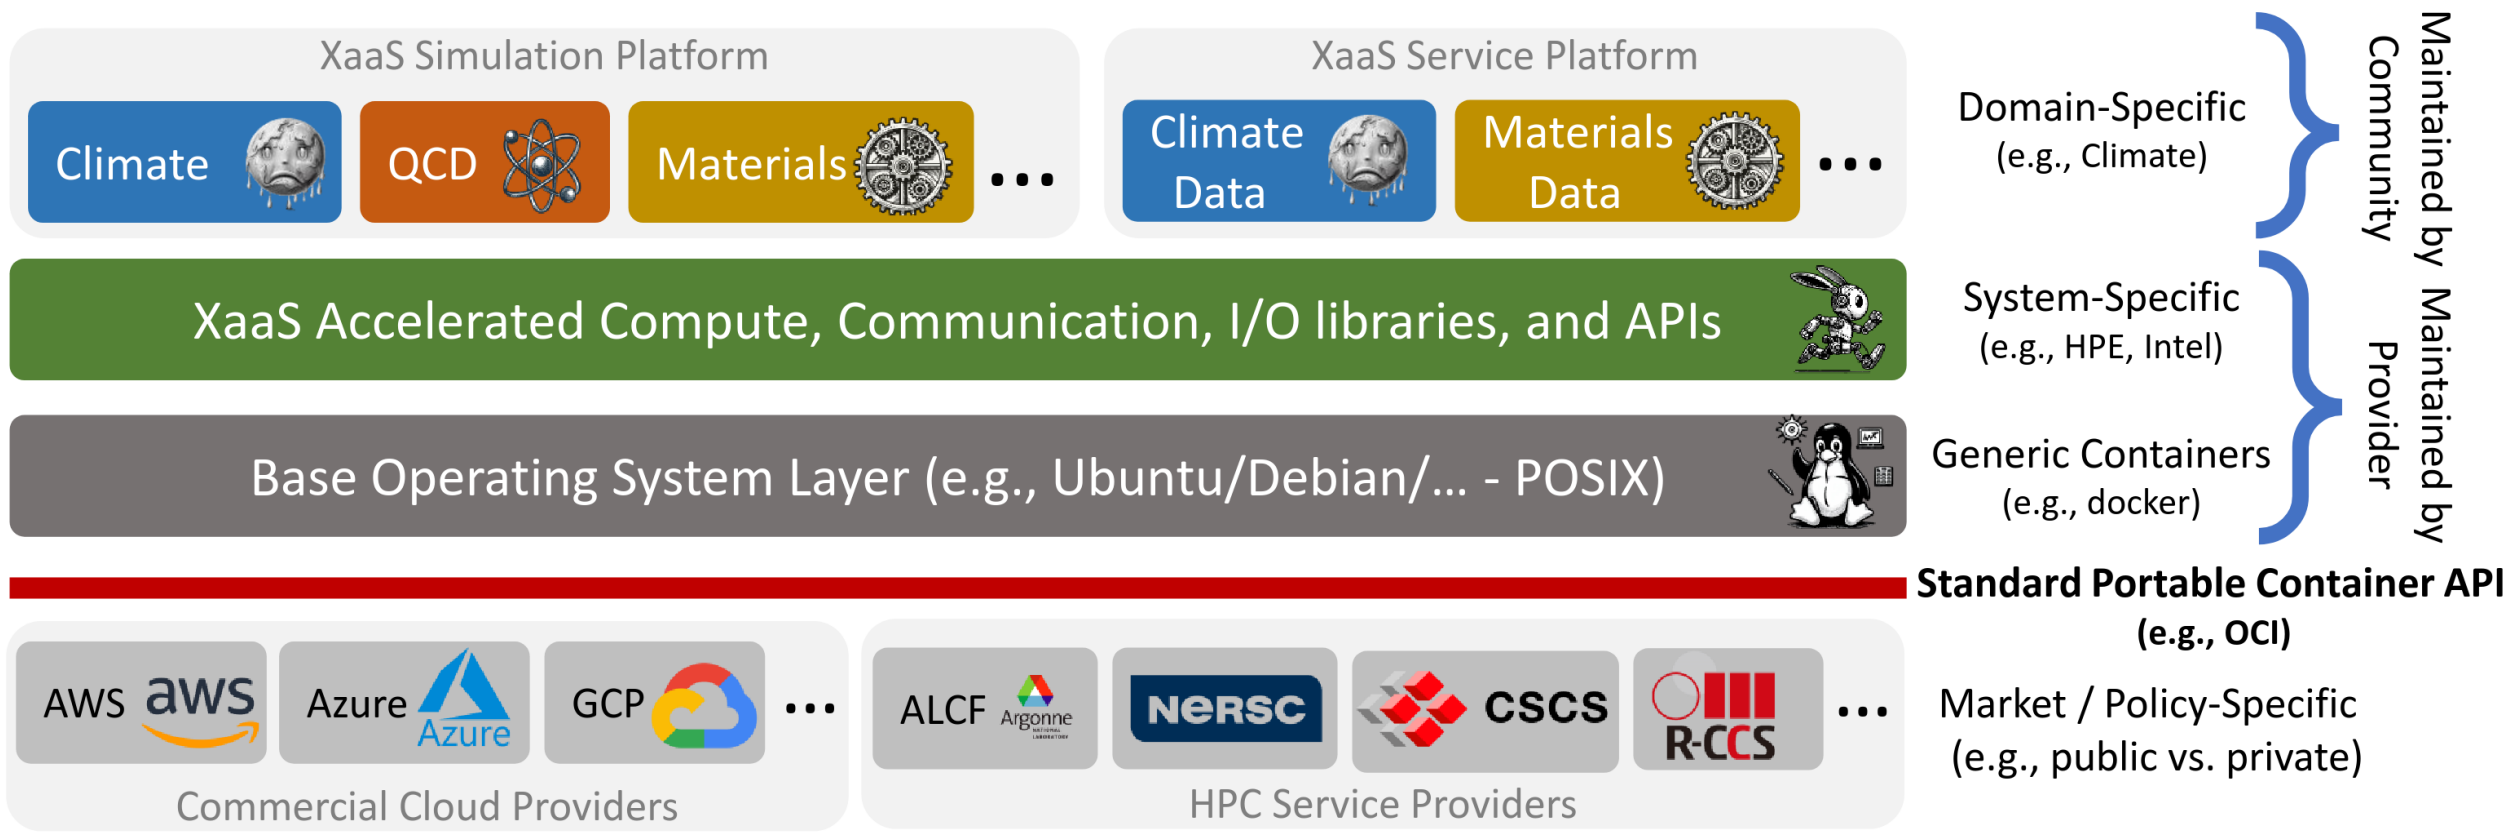
HPC Service Providers

Market / Policy-Specific
(e.g., public vs. private)

XaaS System Architecture



XaaS System Architecture



XaaS in Practice



Spack

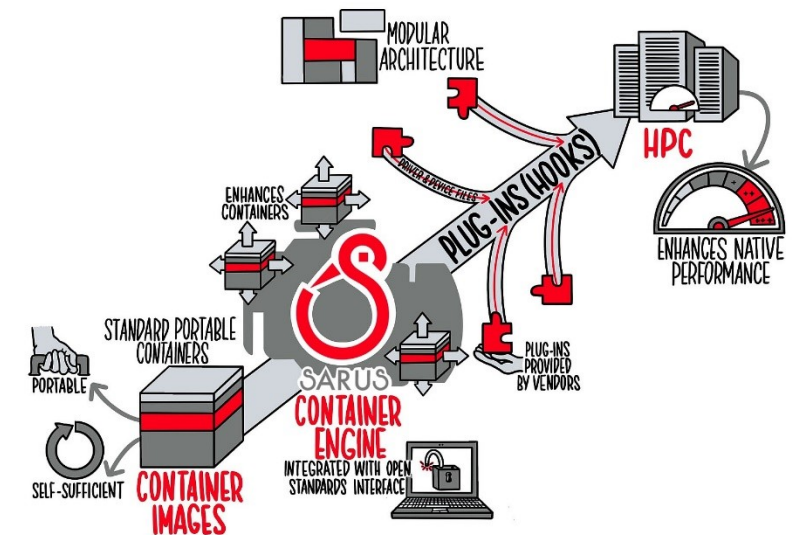
Full Build & Specialization

XaaS in Practice



Spack

Full Build & Specialization




Flexible Library Hooks

XaaS in Practice



Spack

Full Build & Specialization



Flexible Library Hooks

XaaS in Practice




Spack

Full Build & Specialization

More Specialization

More Productivity




Flexible Library Hooks

XaaS in Practice



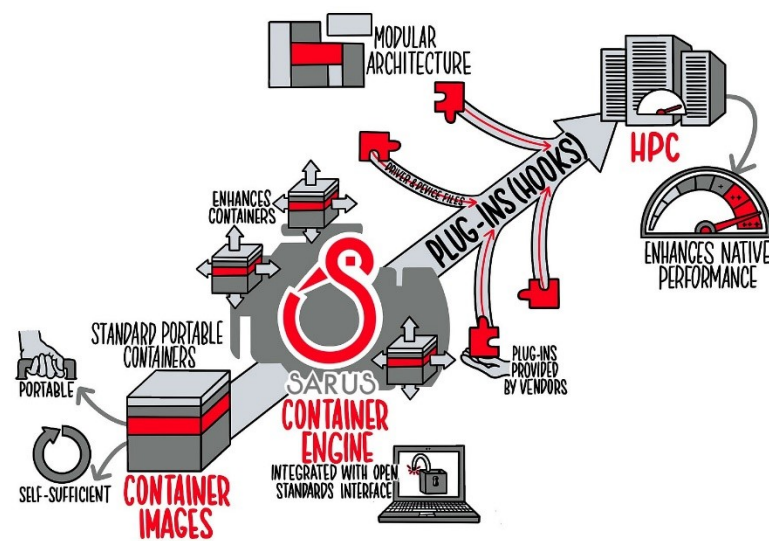

Spack

Full Build & Specialization



aCe

Partial Recompilation



Flexible Library Hooks

What XaaS Can Offer Us?

What XaaS Can Offer Us?

**Cloud
Productivity**

What XaaS Can Offer Us?

**Cloud
Productivity**

**Performance
Portability**

What XaaS Can Offer Us?

**Cloud
Productivity**

**Performance
Portability**

Unified Interface

What XaaS Can Offer Us?

**Cloud
Productivity**

**Performance
Portability**

Unified Interface

**Flexible
Scheduling**




Dan Ernst
NVIDIA



Ian Foster
Argonne National Laboratory



Torsten Hoefler
ETH Zurich



Thomas Schulthess
Swiss National Supercomputing
Centre (CSCS)



Marcin Copik
ETH Zurich

Q&A

